

Deep RL

I Assignment 1


I.1 Flow Matching Policy vs. MSE Policy

The flow policy agent takes actions to achieve the goal in a more orderly and consistent way. In contrast the MSE policy agent either achieves the goal quickly and efficiently or it gets very off track and performs actions that appear not to be leading directly toward pushing the T to the target location. A few hypotheses as to what might be going wrong here:


- the MSE policy isn't as equipped to handle multimodal distributions and when in a scenario where it should pick a mode it instead splits the difference leading to a suboptimal action.
- the MSE policy may be less consistent because its predictions are worse than the flow policy and as a result it gets out of distribution more often.

II Assignment 2

II.1 3.2 Questions



(a) Small Batch Size



(b) Large Batch Size

- Which value estimator has better performance without advantage normalization: the trajectory-centric one, or the one using reward-to-go?
 - Reward to go (RTG) has better performance without advantage normalization. You can see that RTG converges to 200 in both the large and small batch cases whereas the baseline does not in the small batch case.
- Between the two value estimators, why do you think one is generally preferred over the other?
 - The RTG estimator delivers higher avg return and lower stdev return when advantage normalization is not used.

- Did advantage normalization help?
 - Yes, this was the most significant impactful change. Turning advantage normalization on increased the speed of learning and decreased the variance.
- Did the batch size make an impact?
 - Yes, the returns are lower stdev for larger batch size.